


Trees

marcel.mukundi


October 2020

See all drawings as attached figures

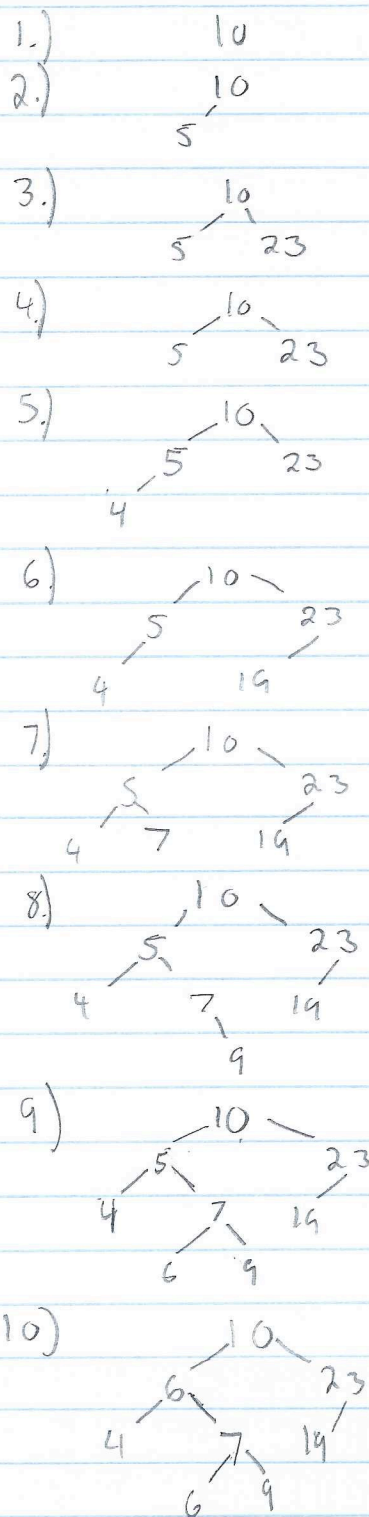
①



②




③




- 4) (a) 5
(b) 4
(c) 2
(d)
 - Preorder: 100, 50, 3, 1, 20, 80, 52, 90, 83, 99, 150, 125, 152
 - Inorder: 1, 3, 20, 50, 52, 80, 83, 90, 99, 100, 125, 150, 152
 - Postorder: 1, 20, 3, 52, 83, 99, 90, 80, 50, 125, 152, 150, 100


⑤




⑥




⑦




⑧




⑨



10A



10B



- 11) Each inner node has M pointers and $M-1$ keys. Each pointer is a 64 bit, or 8 byte, value according to the architecture, and each UUID is 128 bits, or $128/8 = 16$ bytes. This results in a total of

$$\begin{aligned} 8M + 16(M - 1) &= 8M + 16M - 16 \\ &= 24M - 16 \end{aligned}$$

24M-16 bytes. Since M seems to be 5 in the visualization, this would be

$$\begin{aligned} 24 \cdot 5 - 16 &= 120 - 16 \\ &= 104 \end{aligned}$$

104 bytes for each internal node.

Each leaf node has L records and a pointer. Each customer record is

$$\begin{aligned} \text{UUID uuid} &= 128/8 = 16 \text{ bytes} \\ \text{char}[32] \text{ name} &= 32 \cdot 2 = 64 \text{ bytes} \\ \text{int ytd_sales} &= 4 \text{ bytes} \\ \text{Total} &= 84 \text{ bytes} \end{aligned}$$

84 bytes long, and a pointer is still 8 bytes. This results in a total of $84L+8$ bytes. Since L seems to be 5 as well, this makes

$$\begin{aligned} 84 \cdot 5 + 8 &= 420 + 8 \\ &= 428 \end{aligned}$$

428 bytes for each leaf node.

On average, a tree with N records will be about $\log_M(N)$ nodes tall.

For example, a tree with 30,000 records will be about $\log_M(30000)$ nodes tall. If $M=5$, this is a height of $\log_5(30000)$, which is about 6 or 7.

By the same logic, a tree with 2,500,000 records will be about $\log_M(2,500,000)$ nodes tall. If $M=5$, this is a height of $\log_5(2,500,000)$, which is about 9.